

Australian Gas
Production
Analysis

- ***Vompolu Sai Tanuj***

PGP-BABI(June 2019) G1

Table of Contents

1) Project Objective

2) Initial Data Analysis

- a) Invoking the required Packages to be used in R Code
- b) Calling the required dataset into R
- c) Inspection of Time Series Data
- d) Analysing the components of Time Series using Decomposition

3) Splitting of Time Series and finding periodicity

- a) Deseasonalizing the Time Series
- b) Splitting of Time Series into Training and Testing data
- c) Checking the periodicity of the data

4) Using Simple forecast methods

- a) Random Walk Model
- b) Simple Exponential Smoothing
- c) Double Exponential Smoothing
- d) Holt Winter's Model

5) Building Regression models

- a) Checking the time series for stationarity
- b) Converting the series into stationary
- c) Performing Manual ARIMA
- d) Performing Auto ARIMA

- e) **Model Comparison**
 - f) **Box Ljung Test**
 - g) **Making final forecasts**
- 6) Project Conclusion
- 7) Appendix – A (Source Code)

1) Project Objective:

The project is based on the **Monthly Gas Production of Australia**. The data given to us contains the **quantity of gas** that has been **produced every month** for a course of time. We are required to use the same data, build various **time forecast models** and use the **best model** out of them to predict what would be the **quantity of the gas production** that would take place for **12 months** beyond the time period in the given data.

2) Initial Data Analysis:

The required dataset is present in one of the packages in the R called “Forecast”. The dataset can be called onto to the session after importing the package.

a) Invoking the necessary libraries into the R session:

The following libraries have to invoke into the R session before we can begin the analysis. The libraries, if not installed, can be installed by using the function

install.packages() and called onto the R session using **library()**.

Forecast – This library contains the dataset “**Gas**” and several other **functions to be used for model building**

- **tseries** – This library contains various functions which will be **useful for modifying time series**.
- **ggplot2** – This library contains functions to be used for **plotting time series data**.
- **dygraphs** – This library contains functions for **plotting time series object**.
- **xts** – This library contains functions useful for **various plots**.
- **fts** – This library is useful to find the **periodicity**
- **TSA** – This library is useful to build **regression models**.
- **Metrics** – This library contains functions for **model performance**.

b) Calling the required dataset into the R session:

The required dataset ‘**gas**’ is downloaded along with the **forecast** package and can be called into the **session** using its variable name ‘**gas**’. For our **further analysis**, we will call it into **another variable** namely ‘**gasprod**’ **which will be used to call the original dataset when required**. The dataset can be viewed using the function **print()**. We can check the type of dataset using the function **class()**.

```
> ### As the dataset to be used for the analysis is present in the ####  
> ### forecast library, the data can be called directly using ###  
> ### the data name 'gas'. For our analysis, we import it and  
> ### name it 'gasprod' ###  
> gasprod = gas  
.
```

```

> ### viewing the data ###
> print(gasprod)
      Jan  Feb  Mar  Apr  May  Jun  Jul  Aug  Sep  Oct  Nov  Dec
1956 1709 1646 1794 1878 2173 2321 2468 2416 2184 2121 1962 1825
1957 1751 1688 1920 1941 2311 2279 2638 2448 2279 2163 1941 1878
1958 1773 1688 1783 1984 2290 2511 2712 2522 2342 2195 1931 1910
1959 1730 1688 1899 1994 2342 2553 2712 2627 2363 2311 2026 1910
1960 1762 1815 2005 2089 2617 2828 2965 2891 2532 2363 2216 2026
1961 1804 1773 2015 2089 2627 2712 3007 2880 2490 2237 2205 1984
1962 1868 1815 2047 2142 2743 2775 3028 2965 2501 2501 2131 2015
1963 1910 1868 2121 2268 2690 2933 3218 3028 2659 2406 2258 2057
1964 1889 1984 2110 2311 2785 3039 3229 3070 2659 2543 2237 2142
1965 1962 1910 2216 2437 2817 3123 3345 3112 2659 2469 2332 2110
1966 1910 1941 2216 2342 2923 3229 3513 3355 2849 2680 2395 2205
1967 1994 1952 2290 2395 2965 3239 3608 3524 3018 2648 2363 2247
1968 1994 1941 2258 2332 3323 3608 3957 3672 3155 2933 2585 2384
1969 2057 2100 2458 2638 3292 3724 4652 4379 4231 3756 3429 3461
1970 3345 4220 4874 5064 5951 6774 7997 7523 7438 6879 6489 6288
1971 5919 6183 6594 6489 8040 9715 9714 9756 8595 7861 7753 8154
1972 7778 7402 8903 9742 11372 12741 13733 13691 12239 12502 11241 10829
1973 11569 10397 12493 11962 13974 14945 16805 16587 14225 14157 13016 12253
1974 11704 12275 13695 14082 16555 17339 17777 17592 16194 15336 14208 13116
1975 12354 12682 14141 14989 16159 18276 19157 18737 17109 17094 15418 14312
1976 13260 14990 15975 16770 19819 20983 22001 22337 20750 19969 17293 16498
1977 15117 16058 18137 18471 21398 23854 26025 25479 22804 19619 19627 18488
1978 17243 18284 20226 20903 23768 26323 28038 26776 22886 22813 22404 19795
1979 18839 18892 20823 22212 25076 26884 30611 30228 26762 25885 23328 21930
1980 21433 22369 24503 25905 30605 34984 37060 34502 31793 29275 28305 25248
1981 27730 27424 32684 31366 37459 41060 43558 42398 33827 34962 33480 32445
1982 30715 30400 31451 31306 40592 44133 47387 41310 37913 34355 34607 28729
1983 26138 30745 35018 34549 40980 42869 45022 40387 38180 38608 35308 30234
1984 28801 33034 35294 33181 40797 42355 46098 42430 41851 39331 37328 34514
1985 32494 33308 36805 34221 41020 44350 46173 44435 40943 39269 35901 32142
1986 31239 32261 34951 38109 43168 45547 49568 45387 41805 41281 36068 34879
1987 32791 34206 39128 40249 43519 46137 56709 52306 49397 45500 39857 37958
1988 35567 37696 42319 39137 47062 50610 54457 54435 48516 43225 42155 39995

1989 37541 37277 41778 41666 49616 57793 61884 62400 50820 51116 45731 42528
1990 40459 40295 44147 42697 52561 56572 56858 58363 45627 45622 41304 36016
1991 35592 35677 39864 41761 50380 49129 55066 55671 49058 44503 42145 38698
1992 38963 38690 39792 42545 50145 58164 59035 59408 55988 47321 42269 39606
1993 37059 37963 31043 41712 50366 56977 56807 54634 51367 48073 46251 43736
1994 39975 40478 46895 46147 55011 57799 62450 63896 57784 53231 50354 38410
1995 41600 41471 46287 49013 56624 61739 66600 60054

>
> ### checking the class of the imported ###
> class(gasprod)
[1] "ts"
>

```

Inferences:

- We can see from the above results that the data contains values from **January of 1956** to the **August of 1995**.
- The data is already a **time series** object and need not be converted to one for further analysis.

c) Inspection of the Time Series data:

The Time Series data needs to be **inspected** using the **basic functions** before the actual analysis can be started. The inspection can be done using the following functions:

- The **summary()** function can be used to get **basic statistical information** on the **distribution of gas production values** in the data.

```
> ### Inspection of the time series data ###  
> summary(gasprod)  
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   
    1646    2675    16788   21415   38629   66600
```

- The **anyNA()** function can be used to check for **any missing values** in the data.


```
> anyNA(gasprod)  
[1] FALSE
```

- The **findfrequency()** function can be used to check the **frequency of the time series**.

```
> findfrequency(gasprod)  
[1] 12
```


- The **ts.plot()** function can be used to **plot the time series**. The additional arguments like **colour, Title, etc.** can be included in a **list** and must be passed through the **gpars** argument.

```
> ts.plot(gasprod,gpars = list(xlab = "Year",ylab = "Gas Production",  
+                             main = "Australian Gas Production (1956-1995)",  
+                             col = c("Blue")),lwd = 2)
```



- The function **dygraph()** can be used to plot the **time series** as a **step plot**. An additional argument **stepPlot = TRUE** must be passed to obtain a step plot.

```
> stepplot = dygraph(gasprod,main = "Australian Gas Production (1956-1995)",
+                   xlab = "Year",ylab = "Gas Production") %>%dyOptions(stepPlot= TRUE,pointSize = 0,fillGraph = TRUE,
+                   fillAlpha = 0.2)
> stepplot
```



Inferences:

- We can see that **minimum amount of gas produced** was **1646 units**.
- **The Maximum amount of gas produced** was **66600 units**.
- The **median of the data** is **16788 units**.
- The **mean of the data** is at **21415 units**.
- In the given **time series data**, there are no **missing values**.
- The frequency of the **time series** is **12** which indicates it is a **monthly time series**.
- From the plot, we can say that the **lowest value of 1646 units** should have been recorded in the year **1956**.
- The **highest value of 66600 units** should have been recorded around the year **1995**.
- The **time series** doesn't display **any trend** until **1970**.
- The **time series** shows **seasonality** throughout the data.
- The **time series** shows an **upward trend** after **1970**.
- The **effect of seasonality** is **little** until **the year 1970**.
- The **seasonality** changes are very **drastic** after the year **1970**.

d) Analyzing the components of Time Series using Decomposition:

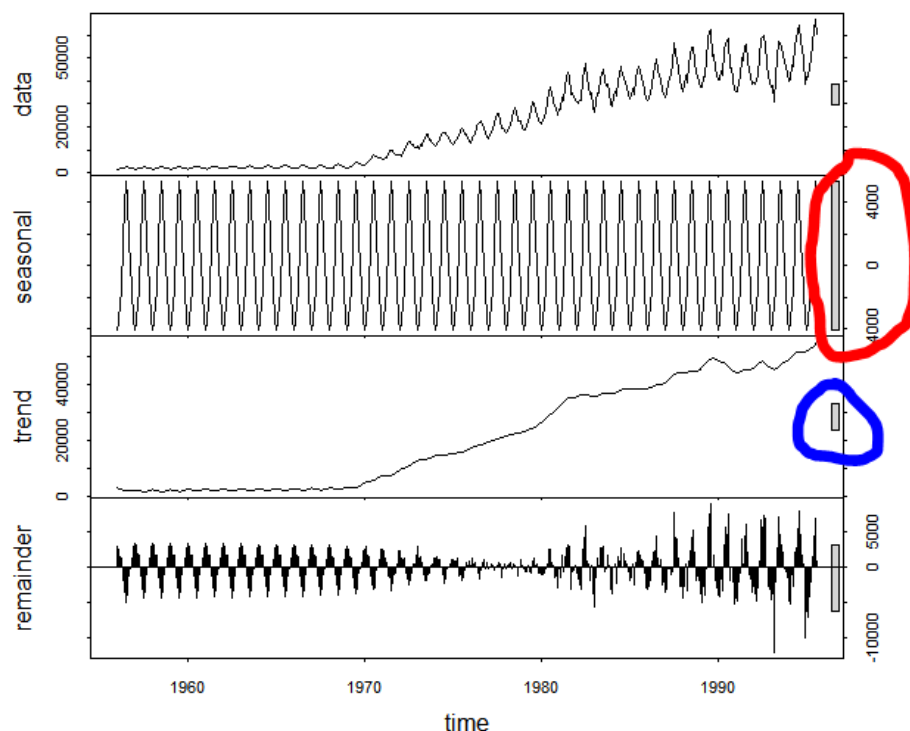
The major **components of any time series data** are **Trend, seasonality and residuals**. These components would be difficult to see in a **normal time series plot**. Therefore a process called **decomposition** can help **analyse each of the components of the time series** by **separating** all the **three components** separately. The **decomposition** can be done

using the function **stl()** with argument **s.window = "periodic"** since the **occurrence of the seasonality is periodic**. After creating a **decomposed object**, the same object can be used to **plot** various plots such as **monthplot** and **seasonalplot**. They can be plotted by using the functions **monthplot()** and **seasonalplot()** respectively.

```

> ### Inspection of individual elements by decomposition of time series ###
> dc.gasprod = stl(gasprod,s.window = "periodic")
> plot(dc.gasprod)

```




```

> monthplot(gasprod,main = "Month Plot for Australian Gas Production")


```

Month Plot for Australian Gas Production



```
> seasonplot(gasprod, year.labels = TRUE,
+           main = "Month Plot for Australian Gas Production")
```

Month Plot for Australian Gas Production



Inferences:

- The **seasonality** seems to be **periodic** almost **resembling** the **same pattern** for the whole time period.
- The **time series** did not **trend** until **1970**, and then after **1970**, it started showing an **upward trend**.
- The **residuals**, which is the remainder after the **trend and seasonality**, is **considerably low** compared to the **trend and seasonality**.
- **Even** though of **higher magnitude**, it actually is **less significant** in describing the **time series**.
- The **magnitude** of **trend** maybe **lower** compared to the **seasonality**, it is **very significant** parameter in determining the **time series data**.
- The **average values** keep increasing **every year** after **1970**.
- The **average values** for the **July** seems to be the highest.
- We can observe **periodic seasonal fluctuations** with an **increasing trend** indicates the **series** follows a **multiplicative seasonality**.


3) Splitting of Time Series and finding periodicity:

a) Deseasonalizing the Time Series:


As inferred above, the time series contains **periodic seasonality** with an **increasing trend**. This is indicative of **multiplicative seasonality** which is high in magnitude. To create certain time series model, the **time series** needs to be eliminated of all the seasonality. To do this, we can use the **time series object** that was created during the **decomposition** and **eliminate** the seasonality component

from it. We name the **deseasonalized time series** as **'ds.gasprod'**.

```
> ### De-Seasonalizing the time series from the decomposed time series####  
> ds.gasprod = (dc.gasprod$time.series[,2]+dc.gasprod$time.series[,3])  
> ts.plot(ds.gasprod,gpars = list(xlab = "Year",ylab = "Gas Production",  
+                               main = "Australian Gas Production (Deseasonlized)",  
+                               col = c("Red")),lwd = 2)  
> |
```




```
> ts.plot(ds.gasprod,gasprod,gpars = list(xlab = "Year",ylab = "Gas Production",  
+                                         main = "Australian Gas Production (Deseasonlized  
+                                         vs. Original)",  
+                                         col = c("Red","Blue")),lwd = 2)  
> |
```




b) Splitting of Time Series into Training and Testing data:

The splitting of the time series data can be done using the **window()** function. We need to split both the **original(Seasonalized)** data and also the **deseasonalized data** into training and testing data. The **Training data** must be considered from **1970 January to 1993 December**. The **Testing data** must be considered from **1994 January to 1995 August**. The Training data for **original and deseasonalized** is named '**train.gasprod**' and '**train.ds.gaspord**'. The Testing data for **original and deseasonalized** is named '**test.gasprod**' and '**test.ds.gaspord**'.


```
> ### Splitting the time series into training and testing samples (Original) ###  
> train.gasprod = window(gasprod,start=c(1970,1), end=c(1993,12), freq=12)  
> ts.plot(train.gasprod)  
,
```




```
> test.gasprod = window(gasprod,start=c(1994,1),end=c(1995,8), freq=12)  
> ts.plot(test.gasprod)
```




```
> autoplot(train.gasprod, series="Train") + autolayer(test.gasprod, series="Test") +
+ ggtitle("Gas Production Training and Test data") +
+ xlab("Year") + ylab("Production") +
+ guides(colour=guide_legend(title="Forecast"))
```




```
> ### splitting the time series into training and testing samples(Deseasonalize) ####
> train.ds.gasprod = window(ds.gasprod,start=c(1970,1), end=c(1993,12), freq=12)
> ts.plot(train.ds.gasprod)
```



```
> test.ds.gasprod = window(ds.gasprod, start=c(1994,1), end=c(1995,8), freq=12)
> ts.plot(test.ds.gasprod)
```



```
> autoplot(train.ds.gasprod, series="Train") + autolayer(test.ds.gasprod, series="Test") +
+ ggtitle("Gas Production Training and Test data(De-seasonalized)") +
+ xlab("Year") + ylab("Production") +
+ guides(colour=guide_legend(title="Forecast"))
```



c) Checking the periodicity of the Time Series:

The **periodicity** of a **time series data** is the measure of **regular intervals** at which the **observations** are recorded. We can check the **periodicity** of the time series data using the function **periodicity()**. We can also check the frequency of the data using **findfrequency()** function.

```
> ### Checking the periodicity of the Time Series ####
> periodicity(gasprod)
Monthly periodicity from Jan 1956 to Aug 1995
> findfrequency(gasprod)
[1] 12
> periodicity(train.gasprod)
Monthly periodicity from Jan 1970 to Dec 1993
> findfrequency(train.gasprod)
[1] 12
> periodicity(test.gasprod)
Monthly periodicity from Jan 1994 to Aug 1995
> findfrequency(test.gasprod)
[1] 12
```

Hence we can see that **periodicity** for **time series data** is **12**.


4) Using Simple Forecasting methods:

The **future values** for the **future** can be forecasted can be done using **various forecast models**. The models are basically of two types. One of them is **Simple Forecast Methods** which use the **pattern of the past values** to predict the **future values**.

a) Random Walk Model:

This method forecasts **next period value** as per the amount of **change over time (called the drift)** is evaluated the **average change** seen in past data. The **function forecast()** on the **decomposed object** along with the argument **method = "rwdrift"**.

```
> ##### Naive Method #####  
> gasprod.rw = stl(train.gasprod,s.window = 'p')  
> gasprod.rw = forecast(de.gasprod.rw,method = "rwdrift",h = 20)  
> ts.plot(test.gasprod,gasprod.rw$mean,gpars = list(col = c("Red","Blue"),  
+         main = "Random walk with Drift(Original vs. Forecasted)"))  
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,  
+       box.lwd = 0.1,cex = 0.75)  
> vec.rw = cbind(test.gasprod,gasprod.rw$mean)  
> MAPE.rw = mean(abs(vec.rw[,1]-vec.rw[,2])/vec.rw[,1])  
> print(MAPE.rw)  
[1] 0.0737419
```



Inferences:

We can see that even though **MAPE** is very low, from the graphs, we can see that the **values forecasted** are always lower than the **Actuals**. We cannot use this **method of**

forecast just because of **MAPE** because there are not many **parameters** that can be **rectified** while making the model and this time series data **needs** the option of **parameter tuning** due to its **complexity**.

b) Simple Exponential Smoothing:

The **Simple Exponential Smoothing** is a **one-step** forecast method where all the **forecast values** are **identical**. The method can be expressed **mathematically** as,

$$Y_{t+1} = \alpha Y_t + \alpha(1-\alpha)Y_{t-1} + \alpha(1-\alpha)^2Y_{t-2} + \dots, \quad 0 < \alpha < 1$$

Where α is the smoothing parameter for the level.

The **simple exponential smoothing** can be done using the function **ses()**

```
> ##### Simple Exponential Smoothing (Original) #####
> gasprod.ses = ses(train.gasprod,start = c(1970,1),end = c(1993,12),frequency = 12,h = 20)
> summary(gasprod.ses)
```

Forecast method: Simple exponential smoothing

Model Information:
Simple exponential smoothing

```
Call:
ses(y = train.gasprod, h = 20, start = c(1970, 1), end = c(1993,
```

```
call:
  12), frequency = 12)
```

```
Smoothing parameters:
  alpha = 0.9999
```

Initial states:
1 = 6002.9039

sigma: 3344.025

AIC	AICc	BIC
6309.126	6309.210	6320.115

Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1
Training set	131.0317	3332.394	2416.021	0.1509165	8.113316	0.9164853	0.3030855

```
> print(gasprod.ses$mean)
```


	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1994	43736.25	43736.25	43736.25	43736.25	43736.25	43736.25	43736.25	43736.25	43736.25	43736.25	43736.25	43736.25
1995	43736.25	43736.25	43736.25	43736.25	43736.25	43736.25	43736.25	43736.25				

```

> ts.plot(test.gasprod,gasprod.ses$mean,gpars = list(col = c("Red","Blue"),
+ main = "Simple Exponential Smoothing(Original vs. Forecasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+ box.lwd = 0.1,cex = 0.75)

```

Simple Exponential Smoothing(Original vs. Forecasted)



```

> vec.ses = cbind(test.gasprod,gasprod.ses$mean)
> MAPE.ses = mape(test.gasprod,gasprod.ses$mean)
> print(MAPE.ses)
[1] 0.1726068

```

```

> #### Simple Exponential Smoothing (Deseasonalized) ####
> gasprod.dses = ses(train.ds.gasprod,start = c(1970,1),end = c(1993,12),frequency = 12,h = 20)
> summary(gasprod.dses)

Forecast method: simple exponential smoothing

Model Information:
simple exponential smoothing

Call:
ses(y = train.ds.gasprod, h = 20, start = c(1970, 1), end = c(1993,
call:
12), frequency = 12)

Smoothing parameters:
alpha = 0.9833

Initial states:
l = 7432.9827

sigma: 2319.347

      AIC      AICC      BIC
6098.373 6098.458 6109.362


Error measures:
      ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 139.963 2311.279 1591.754 0.2300846 6.077039 0.6038107 -0.005553574

> print(gasprod.dses$mean)
      Jan      Feb      Mar      Apr      May      Jun      Jul      Aug      Sep      Oct      Nov
1994 47067.77 47067.77 47067.77 47067.77 47067.77 47067.77 47067.77 47067.77 47067.77 47067.77 47067.77
1995 47067.77 47067.77 47067.77 47067.77 47067.77 47067.77 47067.77 47067.77
      Dec
1994 47067.77
1995

> ts.plot(test.ds.gasprod,gasprod.dses$mean,gpars = list(col = c("Red","Blue"),
+ main = "Simple Exponential Smoothing(Original vs. Forecasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+ box.lwd = 0.1,cex = 0.75)

```

Simple Exponential Smoothing(Original vs. Forecasted)



```
> MAPE.dses = mape(test.ds.gasprod,gasprod.dses$mean)
> print(MAPE.dses)
[1] 0.1110184
```

Inferences:

We can see that **MAPE value(0.11)** is **high** for this model and the graph very much deviates from the actual values. This is because the **time series** data has **both trend and seasonality** which is not suitable for **SES** method.

c) Double Exponential Smoothing:

This method is an **extension** of **simple exponential smoothing**. It contains **two parameters** instead of **one**. Mathematically, it can be expressed in the following way.

$$Y_{t+1} = l_t + b_t$$

Where, l_t and b_t are **estimate of level** and **estimate of trend** respectively

The function **holt()** can be used to build the model

```
> ##### Double Exponential Method (Holt Model) (Original) #####
> gasprod.holt = holt(train.gasprod ,start=c(1970,1),end=c(1993,12), freq=12,h=20)
> summary(gasprod.holt)
```

Forecast method: Holt's method

Model Information:
Holt's method

Call:
holt(y = train.gasprod, h = 20, start = c(1970, 1), end = c(1993,

call:
12), freq = 12)

Smoothing parameters:
alpha = 0.9436
beta = 0.5881

Initial states:
l = 3775.0087
b = 467.2663

sigma: 3542.179

AIC	AICC	BIC
6344.263	6344.476	6362.578

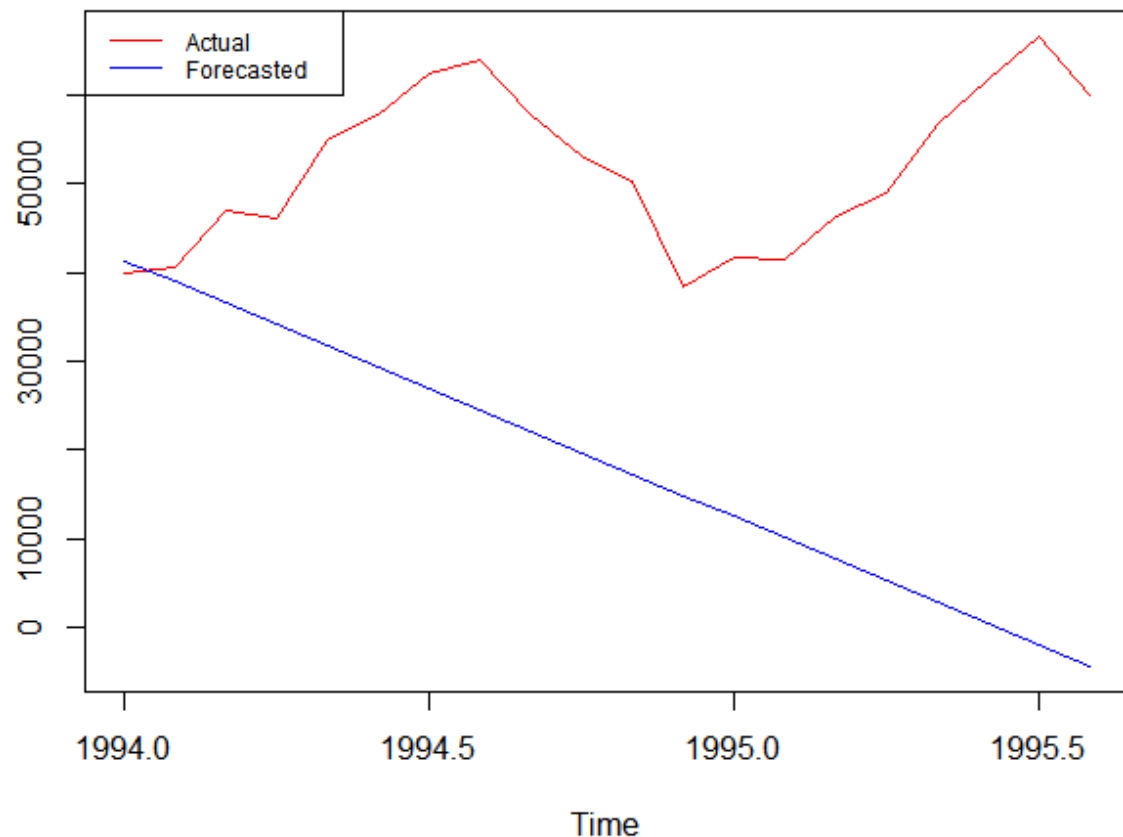
Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1
Training set	-16.95919	3517.494	2436.963	0.2550371	7.909274	0.9244294	-0.01120745

```
> gasprod.holt$mean
      Jan      Feb      Mar      Apr      May      Jun      Jul      Aug      Sep      Oct
1994 41337.3687 38932.0524 36526.7362 34121.4199 31716.1037 29310.7874 26905.4711 24500.1549 22094.8386 19689.5224
1995 12473.5736 10068.2573  7662.9411  5257.6248  2852.3086   446.9923 -1958.3239 -4363.6402
      Nov      Dec
1994 17284.2061 14878.8899
1995

> ts.plot(test.gasprod,gasprod.holt$mean,gpars = list(col = c("Red","Blue"),
+             main = "Holt's Method(Original vs. Forecasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+       box.lwd = 0.1,cex = 0.75)
```

Holt's Method(Original vs. Forecasted)



```
> MAPE.holt = mape(test.gasprod,gasprod.holt$mean)
> print(MAPE.holt)
[1] 0.6201079

> ##### Double Exponential Method (Holt Model) (Deseasonlaize) #####
> gasprod.dholt = holt(train.ds.gasprod,start = c(1970,1),end = c(1993,12),freq = 12,h = 20)
> summary(gasprod.dholt)

Forecast method: Holt's method

Model Information:
Holt's method

call:
holt(y = train.ds.gasprod, h = 20, start = c(1970, 1), end = c(1993,
12), freq = 12)

Smoothing parameters:
alpha = 0.9809
beta = 1e-04

Initial states:
l = 7297.2832
b = 70.1218


sigma: 2324.336

      AIC      AICC      BIC
6101.590 6101.803 6119.905
```

```

> gasprod.dholt$mean
      Jan      Feb      Mar      Apr      May      Jun      Jul      Aug      Sep      Oct      Nov
1994 47142.87 47214.95 47287.03 47359.11 47431.19 47503.27 47575.35 47647.43 47719.51 47791.59 47863.67
1995 48007.83 48079.91 48151.99 48224.07 48296.15 48368.23 48440.31 48512.39
      Dec
1994 47935.75
1995
> ts.plot(test.ds.gasprod,gasprod.dholt$mean,gpars = list(col = c("Red","Blue"),
+               main = "Holt's Method(Original vs. Forecasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+       box.lwd = 0.1,cex = 0.75)

```



```

> MAPE.dholt = mape(test.ds.gasprod,gasprod.dholt$mean)
> print(MAPE.dholt)
[1] 0.1033223

```

Inferences:

We can see that **MAPE value(0.10)** is **high** for this model and the graph very much deviates from the actual values. This is because the **time series** data has **both trend and seasonality** which is not suitable for **Holt's** method.

d) Holt Winter's Model:

This model is an **extension** of the **Holt's model** where instead of **two**, we consider **three parameters**. Mathematically, this can be represented as,

Forecast equation: $Y_{t+1} = l_t + b_t + s_{t-m}(k+1)$

Level Equation: $l_t = \alpha(Y_t - s_{t-m}) + \alpha(1-\alpha)Y_{t-1}$, $0 < \alpha < 1$

Trend Equation: $b_t = \beta(l_t - l_{t-1}) + (1-\beta)b_{t-1}$, $0 < \beta < 1$

Seasonal Equation: $s_t = \gamma(Y_t - l_t - 1 - b_{t-1}) + (1-\gamma)s_{t-m}$, $0 < \gamma < 1$

Here, α , β and γ are **smoothing parameters**.

The **Holt Winter's model** can be applied to the **time series data** using the function **hw()**.

```
> #### Holt winter's method (Original) ####
> gasprod.hw = hw(train.gasprod, start = c(1970,1), end = c(1993,12), freq = 12, h = 20)
> summary(gasprod.hw)

Forecast method: Holt-Winters' additive method

Model Information:
Holt-Winters' additive method

Call:
hw(y = train.gasprod, h = 20, start = c(1970, 1), end = c(1993,
12), freq = 12)

Smoothing parameters:
alpha = 0.3408
beta  = 1e-04
gamma = 0.5936

Initial states:
l = 6253.203
b = 119.5506
s = -4511.742 -2141.073 234.0438 2010.202 5919.329 7284.465
    5272.426 2485.985 -2602.642 -3068.389 -5131.983 -5750.62

sigma: 2109.356

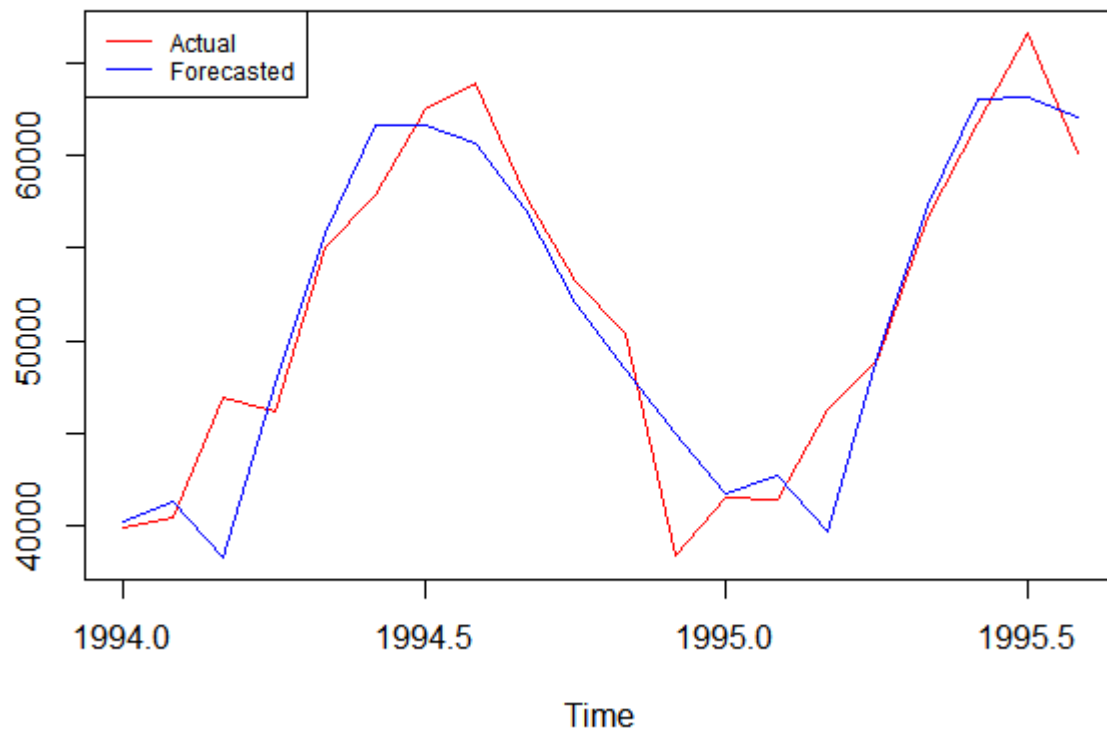
      AIC      AICC      BIC
6057.255 6059.521 6119.525

Error measures:
      ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 78.69136 2049.926 1551.842 0.4623734 7.505829 0.5886704 0.2738151
```

```
> gasprod.hw$mean
      Jan      Feb      Mar      Apr      May      Jun      Jul      Aug      Sep      Oct      Nov
1994 40279.65 41289.25 38304.76 47712.40 55810.29 61602.40 61619.91 60595.61 57042.03 51962.18 48460.99
1995 41741.45 42751.05 39766.56 49174.20 57272.09 63064.20 63081.72 62057.41
      Dec
1994 44992.38
1995
```

```
> ts.plot(test.gasprod,gasprod.hw$mean,gpars = list(col = c("Red","Blue"),
+                                                    main = "Holt winter's(Original vs. Forecasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+       box.lwd = 0.1,cex = 0.75)
```

Holt Winter's(Original vs. Forecasted)



```
> MAPE.hw = mape(test.gasprod,gasprod.hw$mean)
> print(MAPE.hw)
[1] 0.04666037
```

```

> ##### Holt winter's method (Deseasonalize) #####
> gasprod.dhw = hw(train.ds.gasprod,start = c(1970,1),end = c(1993,12),freq = 12,h = 20)
> summary(gasprod.dhw)

Forecast method: Holt-winters' additive method

Model Information:
Holt-winters' additive method

Call:
hw(y = train.ds.gasprod, h = 20, start = c(1970, 1), end = c(1993,
12), freq = 12)

Smoothing parameters:
alpha = 0.3408
beta = 1e-04
gamma = 0.5934

Initial states:
l = 6263.4243
b = 119.8431
s = -1190.992 -669.5463 69.8913 481.8762 1663.536 1972.405
1680.539 771.3246 -793.5156 -903.362 -1416.127 -1666.028

sigma: 2109.394


      AIC      AICC      BIC
6057.265 6059.531 6119.535

Error measures:
      ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 77.89129 2049.962 1551.948 -1.123422 8.141496 0.5887106 0.2739442

> gasprod.dhw$mean
      Jan      Feb      Mar      Apr      May      Jun      Jul      Aug      Sep      Oct      Nov      Dec
1994 44364.14 45005.85 40472.48 49522.52 54097.59 58011.54 56309.27 56342.08 55515.56 51799.98 49934.90 48315.63
1995 45829.18 46470.89 41937.52 50987.55 55562.62 59476.58 57774.30 57807.12
> ts.plot(test.ds.gasprod,gasprod.dhw$mean,gpars = list(col = c("Red","Blue"),
+ main = "Holt winter's(Original vs. Forecasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+ box.lwd = 0.1,cex = 0.75)

```

Holt Winter's(Original vs. Forecasted)



```
> MAPE.dhw = mape(test.ds.gasprod,gasprod.dhw$mean)
> print(MAPE.dhw)
[1] 0.0458347
```

Inferences:


We can see that **MAPE value(0.04)** is **very low** for this model and the graph shows that **the forecasted values** almost **explain the fluctuations** that the **time series data**.

Even the **AIC and BIC values** are **good enough**.

Out of these both models where we used **original** or **deseasonalized**, we can go for **original** since there isn't much difference in **MAPE** between models **created** from **original** data and **deseasonalized** data. And also if we chose **deseasonalized model**, we are giving up **some of the original** data that was lost during the **deseasonalization**. Therefore choosing the **original model** would be the right decision over the **deseasonalized model**.

5) Building Regression models:

The concept of Regression can also be applied to the **time series data** and can be used for **forecasting**. The **forecasted values** can be termed as **Response variable** and the previous values can be termed as **Regressor variables**. But in this type of regression, instead of **correlation**, we talk about **auto-correlation** which showcases the **relation between the time series and the lagged version of itself over several time periods**. The most common regression model that can be built for this time series is **ARIMA**. There are **three components** in the **ARIMA model**. The **differencing factor(d)**, the **Moving average factor(q)** and the **autoregressive term (p)**. In R



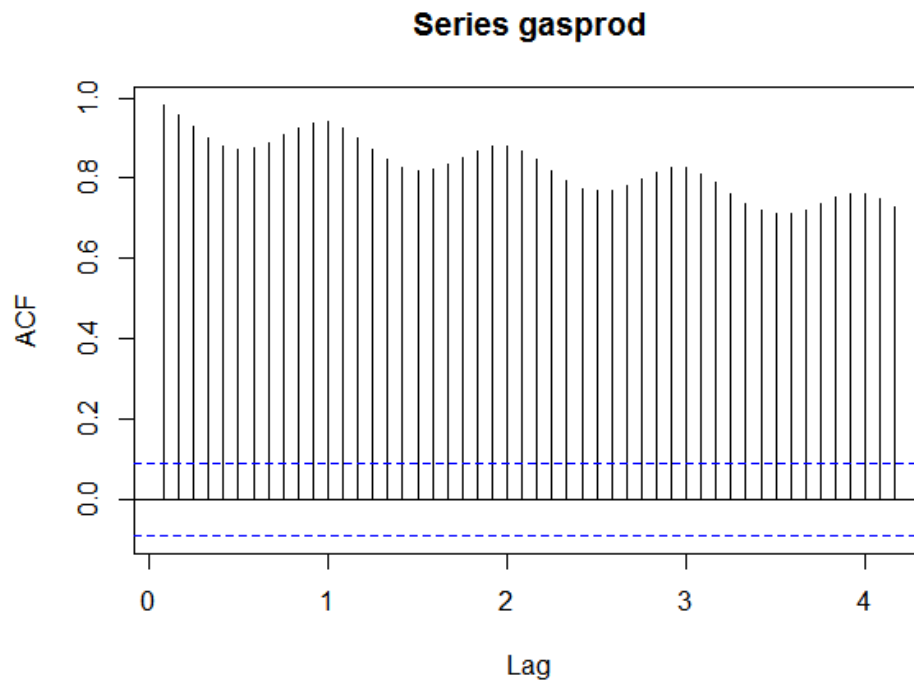
We can see that **the series** has **seasonality and trend**. Therefore by this we can rule out that the series is **non-stationary**.

➤ **Checking the ACF and PACF plots:**

ACF(Auto Correlation Function) and **PACF(Partial Auto Correlation Function)** plots are **two** important plots which show to what extent the **time series** are auto correlated. The plot shows to how many number of **lags**, the correlation is **significant**. The function **acf()** and **pacf()** can be used to get the plots.

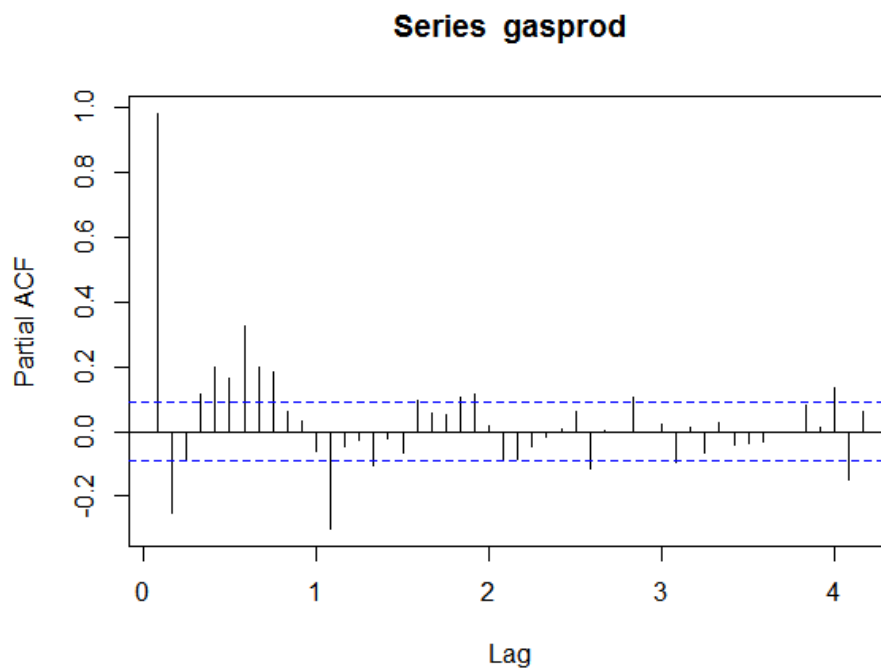
ACF plot:

```
> acf(gasprod, lag.max = 50)
```



PACF plot:

```
> pacf(gasprod, lag.max = 50)
```



From the **ACF** plot, we can see that the values of correlation have not dropped to **zero** even after **50 lags** which is indicative of **non-stationary series**.

➤ **Augmented Dickey-Fuller Test:**

It is a formal test to check whether a **time series** is **stationary or non-stationary**.

H_0 : Time series is non-stationary

H_1 : Time series is stationary

The function **adf()** can be used to perform the test. The **p-value** of the test determines whether the series is stationary or not.

```
> ##### Checking for stationarity using Augmented Dicky- Fuller Test #####  
> adf = adf.test(gasprod, alternative = "stationary")  
> adf
```

Augmented Dickey-Fuller Test

```
data: gasprod  
Dickey-Fuller = -2.7131, Lag order = 7, p-value = 0.2764  
alternative hypothesis: stationary
```

```
> print(adf$alternative)  
[1] "stationary"  
> if(adf$p.value < 0.05){  
+   print("The series is Stationary & Null Hypothesis is rejected")  
+ } else{  
+   print("The series is not Stationary & Null Hypothesis is accepted")  
+ }  
[1] "The series is not Stationary & Null Hypothesis is accepted"
```


As we can see from the above results, the **p-value** is **0.27** and we **reject the alternative hypothesis** that the series is stationary.

b) Converting the series into stationary:


The time series data needs to be converted to **stationary** in order to perform a **Manual ARIMA**. The **series** can be converted to **stationary series** by the process of **differencing**. It is the process in which a new series is got by computing the differences between consecutive observations. This can be done in R with the help of the function **diff()** with the **argument differences = 1** since from the **PACF and ACF** plots, it is evident the differencing value should be taken as **1**.

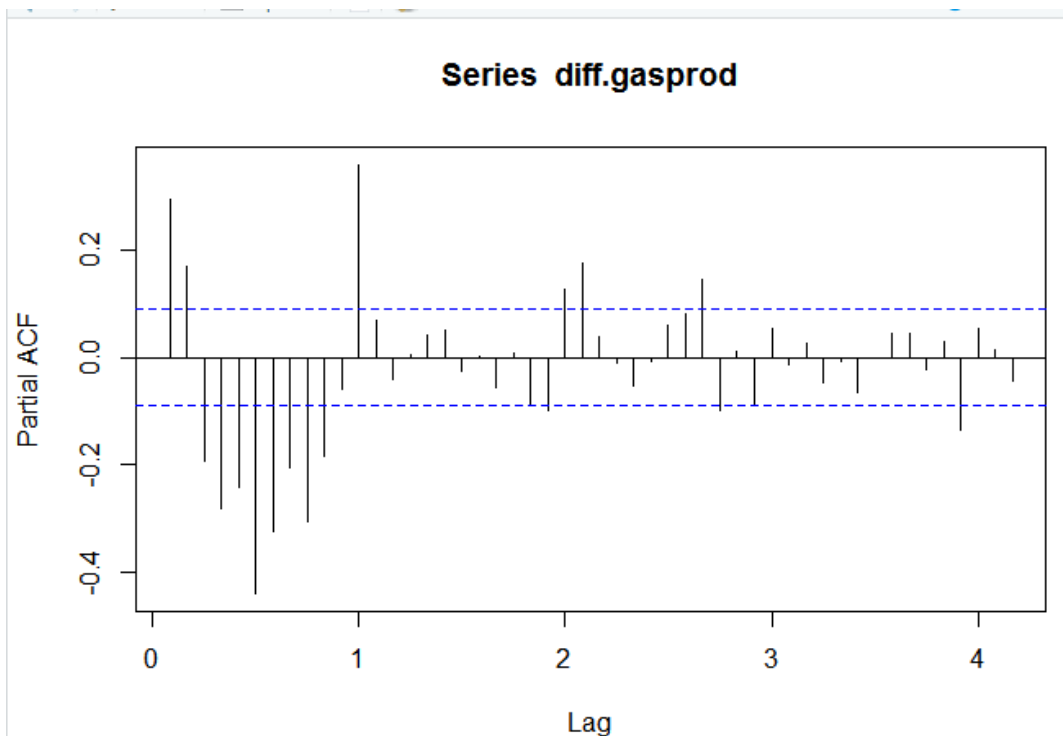

```
#### Stationarizing the series for performing Manual Arima ####
diff.gasprod = diff(gasprod, differences = 1)
ts.plot(diff.gasprod,col = c("Blue"),lwd = 1,main = "Differenced Series")
abline(a=1,b=0,col = c("Red"),lwd =4)
```

Differenced Series



Series diff.gasprod






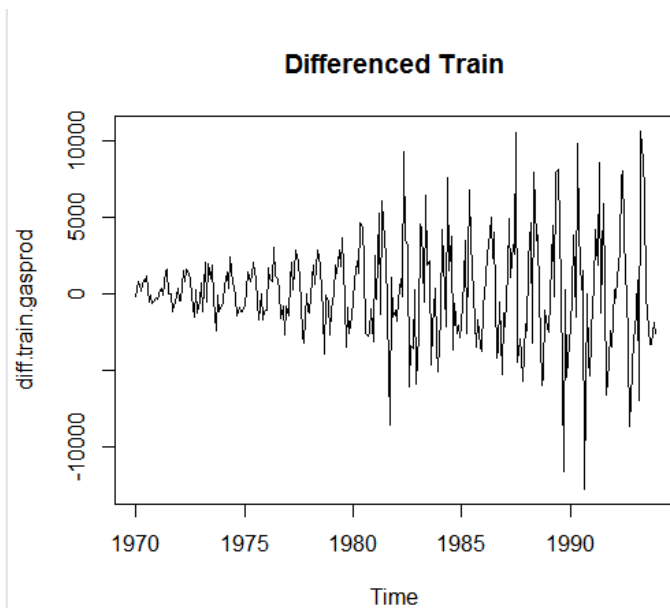
```
> adf.d = adf.test(diff.gasprod, alternative = "stationary")
warning message:
In adf.test(diff.gasprod, alternative = "stationary") :
  p-value smaller than printed p-value
> adf.d
```

Augmented Dickey-Fuller Test

```
data: diff.gasprod
Dickey-Fuller = -19.321, Lag order = 7, p-value = 0.01
alternative hypothesis: stationary

> print(adf.d$alternative)
[1] "stationary"
> if(adf.d$p.value < 0.05){
+   print("The series is Stationary & Null Hypothesis is rejected")
+ } else{
+   print("The series is not Stationary & Null Hypothesis is accepted")
+ }
[1] "The series is Stationary & Null Hypothesis is rejected"

> ### Splitting of time series into Training and Testing Data ###
> diff.train.gasprod = window(diff.gasprod, start = c(1970,1), end = c(1993,12), frequency = 12)
> diff.test.gasprod = window(diff.gasprod, start = c(1994,1), end = c(1995,8), frequency = 12)
> plot.ts(diff.train.gasprod, main = "Differenced Train")
> plot.ts(diff.test.gasprod, main = "Differenced Test")
```



c) Performing Manual ARIMA:

The **Manual ARIMA** requires three parameters namely the **differencing factor(d)**, the **Moving average factor(q)** and the **autoregressive term (p)**. The **d** term has already been calculated during **differencing** which is **1**. From the **PACF** and **ACF**, the **p** term and **q** term are **2** and **2** respectively since the lag values in **ACF** are significant only till **2**. We can use the function **arima()** with the arguments **(2,1,2)**.

```


> marima.gasprod = Arima(diff.train.gasprod,order = c(2,1,2))
> marima.gasprod
Series: diff.train.gasprod
ARIMA(2,1,2)

Coefficients:
          ar1      ar2      ma1      ma2
      -0.0131  0.2545 -0.7212 -0.2788
s.e.    0.1623  0.0695  0.1595  0.1593

sigma^2 estimated as 9895599: log likelihood=-2719.06
AIC=5448.12  AICc=5448.33  BIC=5466.41
> forc.marima.gasprod = forecast(marima.gasprod,h = 20)
> ts.plot(diff.test.gasprod,forc.marima.gasprod$mean,gpars = list(col = c("Red","Blue"),
+ main = "Manual Arima(Original vs. Forecaste
d)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+ box.lwd = 0.1,cex = 0.75)

```

Manual Arima(Original vs. Forecasted)



```

> vec.marima = cbind(diff.test.gasprod,forc.marima.gasprod$mean)
> MAPE.marima = mean(abs(vec.marima[,1]-vec.marima[,2])/vec.marima[,1])
> print(MAPE.marima)
[1] 0.1915037

```

Inferences:

We can see that the **MAPE(0.19)** is **very high** and **AIC** and **BIC** values are on the **higher side**. The **plot** shows that **forecasted**

values and **actual values** don't match. This is due to **multiplicative seasonality** that is present in the dataset.

d) Performing Auto ARIMA:

In the **Auto ARIMA**, we need not convert the series into **stationery series** and also we need not determine the values of **p, d and q**. The function **auto.arima()** with argument **seasonal = TRUE** can be used to perform **Auto ARIMA**.

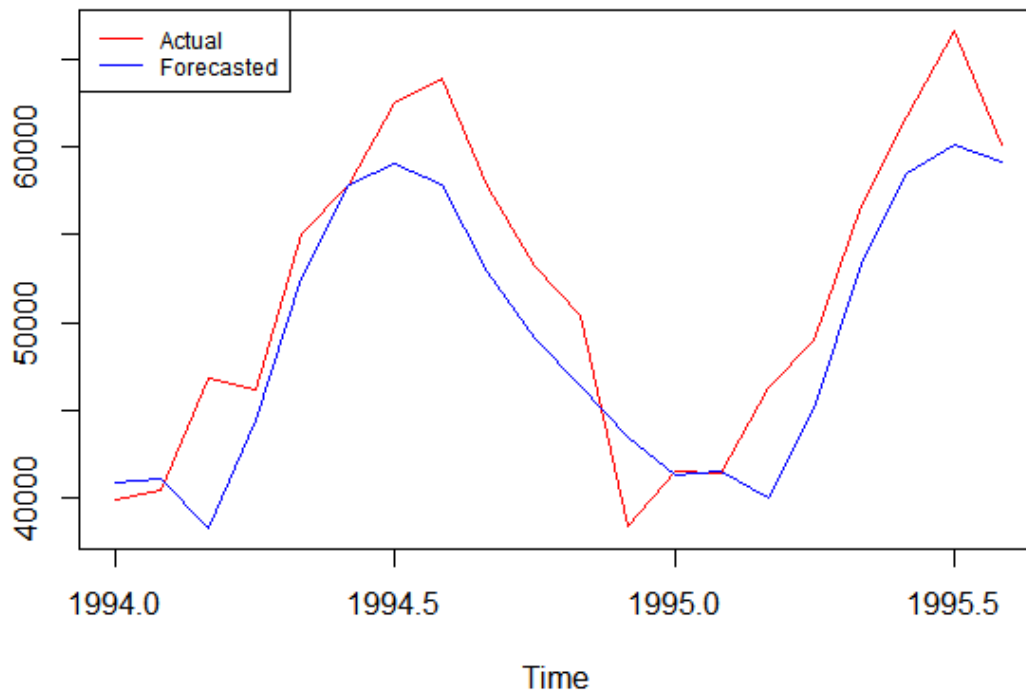
```
> ###USING AUTO ARIMA (Original) ####
> aarima.gasprod = auto.arima(train.gasprod,seasonal = TRUE)
> aarima.gasprod
Series: train.gasprod
ARIMA(2,1,1)(0,1,2)[12]

Coefficients:
          ar1      ar2      ma1      sma1      sma2
      0.5017  0.2057 -0.9583 -0.4404 -0.1236
s.e.  0.0738  0.0722  0.0426  0.0676  0.0639

sigma^2 estimated as 3535010:  log likelihood=-2463.67
AIC=4939.33  AICc=4939.64  BIC=4961.03
> forc.aarima = forecast(aarima.gasprod,h = 20)

> forc.aarima = forecast(aarima.gasprod,h = 20)
> ts.plot(test.gasprod,forc.aarima$mean,gpars = list(col = c("Red","Blue"),
+                                                     main = "Auto Arima(Original vs. Fo
recasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+       box.lwd = 0.1,cex = 0.75)
> |
```

Auto Arima(Original vs. Forecasted)



```
> vec.aarima = cbind(test.gasprod,forc.aarima$mean)
> MAPE.aarima = mean(abs(vec.aarima[,1]-vec.aarima[,2])/vec.aarima[,1])
> print(MAPE.aarima)
[1] 0.06370687
```

Inferences:

We can see that the **MAPE(0.06)** is **very low** and the **plot** shows that **forecasted values** and **actual values** actually match. Even the **AIC** and **BIC** values are **pretty low**. This is due to the reason being the **auto.arima()** using the **best iterations of p,d and q** values to determine the best model.

e) Model Comparison:

For the **time series** data, the following **model measures** are to be kept in mind for comparison:

- **AIC – Akaike Information Criterion** is an important **model measure** which actually helps in model comparison rather than giving us the model performance as it is. **AIC** is the

measure of how much information is being lost by the model while it tries to explain the process. **The lesser the value, the better the model is in prediction.**

- **BIC – Bayesian Information Criterion** is an important **model comparison measure**. It is the **posterior probability in the Bayesian setup**, which gives the **likelihood of the model to the true model**. Therefore, **lesser the value, better the model is in likelihood**.
- **MAPE – Mean Absolute Percentage Error** is the measure of the prediction accuracy of a forecasting method which is represented in the form of **percentage**. **The lesser the value, better the model that fits.**

We have created several **forecasting methods** to **forecast values** using the **training data, original and deseasonalized**, to forecast values and compare it with **testing data, original and deseasonalized**. We have also created several **regression models** using the **differenced data** for **Manual ARIMA** and **original data** for **Auto ARIMA**.

Model	AIC	BIC	MAPE
Random Walk with Drift	NA	NA	7%
Simple Exponential Smoothing - Original Data	6309	6320	17%
Simple Exponential Smoothing - Deseasonalized Data	6098	6109	11%
Double Exponential Smoothing - Original Data	6344	6362	62%
Double Exponential Smoothing - Deseasonalized Data	6101	6119	10%
Holt-Winter's Method - Original Data	6057	6119	4.60%
Holt-Winter's Method - Deseasonalized Data	6057	6119	4.50%
Manual ARIMA - Differenced Data	5448	5466	19%
Auto ARIMA - Original Data	4939	4961	6%

Inferences:

We can see that out of all these models, the **Holt-Winter's Model** with **Deseasonalized data** has the **lowest MAPE value of 4.5%**. But comparing the **Holt-Winter's Model** with **Original data**, it isn't much difference, therefore we can consider **the model with original data over the deseasonalized one**.

The lowest AIC and BIC values belong to the **Auto ARIMA model** with **4939 and 4961 respectively**.

So we are now faced with the question of whether to choose whether **Holt-Winter's model** or **Auto ARIMA**.

<u>Model</u>	<u>AIC</u>	<u>BIC</u>	<u>MAPE</u>
Holt-Winter's Method - Original Data	6057	6119	4.60%
Auto ARIMA - Original Data	4939	4961	6%

If we compare both of them side-by-side, we can see that for a **small change in MAPE value (1.4%)**, we are compromising a **huge amount of AIC and BIC values respectively (around 1000)**. And also we need to consider that real low values of **MAPE** can cause **over fitting**. While **Holt-Winter's model** take only trend and seasonality into account, **ARIMA** considers Moving Average, Difference Term, Auto Regressive Term and also Error term which is an important aspect since as we saw in the decomposition plot that **error term** had a significant impact on the time series data.

Keeping in mind the above reasons, we consider Auto ARIMA model to be the best model to be used for forecasting.

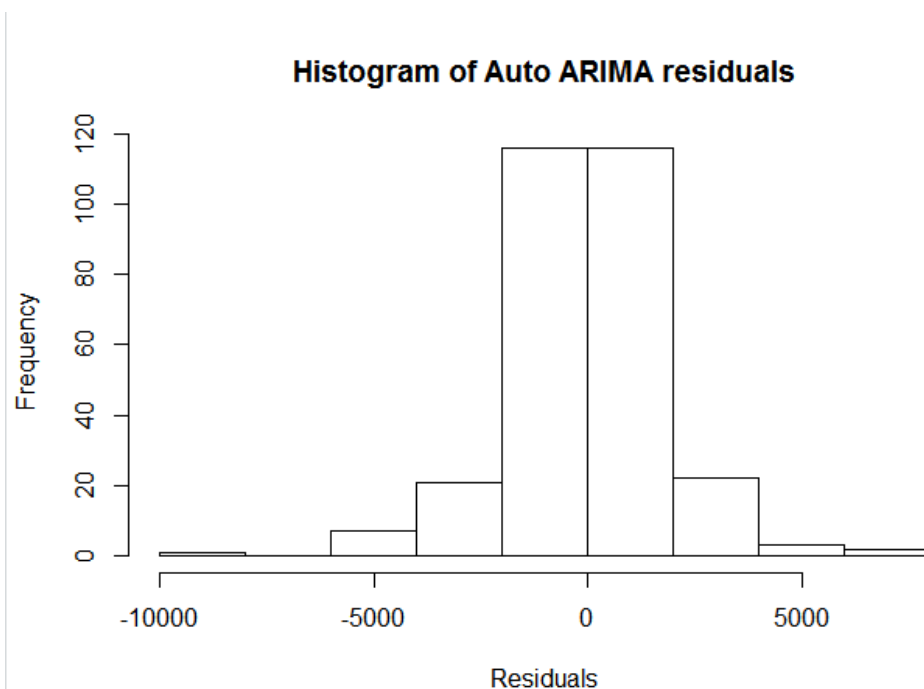
f) Box Ljung Test:

The **Box Ljung Test** is used to check whether the **residuals are following the white noise or not**. It checks whether the residuals are stationary or not.

H0: Residuals are stationary

H1: Residuals are not stationary

```
> ### Ljung box test ###  
> Box.test(aarima.gasprod$residuals,type = "Ljung-Box")  
  
Box-Ljung test  
  
data: aarima.gasprod$residuals  
X-squared = 0.009919, df = 1, p-value = 0.9207  
  
> hist(aarima.gasprod$residuals,main = "Histogram of Auto ARIMA residuals",  
+       xlab = "Residuals")
```



From the above results, we can see that the **P-Value** is more than **0.05** and hence we **reject null hypothesis** and conclude that the **residuals are stationary** concluding they are **independent**. We can also see that the **residuals** follow a


normal distribution (Bell Curve) meaning the residuals are stationary.

g) **Making final forecasts:**

As we have concluded that **Auto ARIMA** gives the best model, we use it to create two models, the one with data ranging from **1970 January to 1993 December (Original Training Data)** and **1970 January to 1995 August (Original Training Data and Testing Data)**. We make the forecast for next 12 time periods i.e. **1995 September to 1996 August**.

```
> ### Making the future forecast using the best model ###
> future = forecast(aarima.gasprod,h = 32)
> plot(future)
> future$mean
```


	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug
1994	40911.17	41136.45	38318.03	44418.69	52575.33	57827.42	59054.39	57800.39
1995	41314.97	41577.63	40068.39	45141.64	53284.34	58496.89	60063.56	59142.29
1996	41882.79	42199.58	40744.41	45855.92	54028.93	59264.56	60849.03	59941.44
	Sep	Oct	Nov	Dec				
1994	52923.33	49148.06	46382.99	43491.75				
1995	54005.17	49942.73	46875.98	43928.71				
1996								



```
> ### Making the future forecast using the model created from gasprod ###
> aarima.fgasprod = auto.arima(gasprod,seasonal = TRUE)
> fgasprod.forc = forecast(aarima.fgasprod,h = 12)
> plot(fgasprod.forc)
> fgasprod.forc$mean
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov
1995									56907.83	52476.28	49719.79
1996	43318.41	43601.71	46668.11	49376.36	57536.25	62184.69	65795.74	63391.54			
	Dec										
1995	43473.70										
1996											

Forecasts from ARIMA(2,1,1)(0,1,1)[12]



6) Project Conclusion:

We were given a **time series data** of the **monthly Australian Gas Production** and were asked to **build a model to best forecast the future values beyond the time periods that data contained**. Various methods such as **decomposition** were used to analyse every component of the time series of data and understand the data better before we go for building a model. Then we applied **deseasonalization** method on the time series and **split the data into training and testing data for both original and seasonalized data**.

After using many forecasting methods and building regression models, we came to conclude that **Auto ARIMA** gave us the **best model** when compared with **other** models using **model measures such as AIC, BIC and MAPE**. Then using this method, we presented the **forecasts** for the **next 12 time periods**.

As we performed the analysis, following things were noted,

- There was **major change** in the **time series data** from the **year 1970**. An **upward trend** started from that time period onwards. This could be due to the **energy crisis of 1970** where the **prices of all the energy sources went up** and the **countries were forced to put more effort into increasing the production of energy sources**.
- The **upward trend** from the year **1970** could also be due to the **technology advancements** that might have occurred in the year **1970** helping in more production of gas.
- There has never been a **decrease** in the **production value** of **gas** from **its preceding year**.
- The **highest value** is in the **year 1995** which was **66600 units** whereas **the lowest value** was in the **year 1956** which was **1646 units**. Even though it has been **39 years**, the **production values** increased by **margin of only 2.4%**
- We can see that there is a **seasonal increase** in the **months of July** which are the **coldest winter months** in **Australia**. The production **increases** in **winter** because **people require more natural gas to fight the cold weather in winters as compared to other months in other seasons**. Hence to meet the **demand**, the production of the natural gas **increases**.

The following **suggestions** can be provided based on the analysis of the data:

- The data collected could have been a **daily production data** as compared to **monthly data** to make **better forecasts**.
- The data had a **major incident affecting** the **analysis** and also the **model building** which could have been avoided by taking a **different time period**.
- It was nowhere mentioned that whether the **monthly production values** taken were either **average** of every day **gas production** or **total of 30 days gas production**. A better analysis could have been done had this **information** been known.

7) Appendix – A (Source Code):

AUSTRALIAN GAS PRODUCTION

ANALYSIS OF AUSTRALIAN GAS PRODUCTION

```
### Invoking of the necessary libraries ###
install.packages('forecast')
library(forecast)
install.packages('tseries')
library(tseries)
install.packages('ggplot2')
library(ggplot2)
install.packages('dygraphs')
library(dygraphs)
install.packages('xts')
library(xts)
install.packages('fts')
library(fts)
install.packages('TSA')
library(TSA)
install.packages('Metrics')
library(Metrics)
> ### As the dataset to be used for the analysis is present in the ###
> ### forecast library, the data can be called directly using ###
> ### the data name 'gas'. For our analysis, we import it and
> ### name it 'gasprod' ###
> gasprod = gas
> ### Viewing the data ###
> print(gasprod)
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	D
ec 1956 25	1709	1646	1794	1878	2173	2321	2468	2416	2184	2121	1962	18
1957 78	1751	1688	1920	1941	2311	2279	2638	2448	2279	2163	1941	18
1958 10	1773	1688	1783	1984	2290	2511	2712	2522	2342	2195	1931	19
1959 10	1730	1688	1899	1994	2342	2553	2712	2627	2363	2311	2026	19
1960 26	1762	1815	2005	2089	2617	2828	2965	2891	2532	2363	2216	20
1961 84	1804	1773	2015	2089	2627	2712	3007	2880	2490	2237	2205	19
1962 15	1868	1815	2047	2142	2743	2775	3028	2965	2501	2501	2131	20
1963 57	1910	1868	2121	2268	2690	2933	3218	3028	2659	2406	2258	20
1964 42	1889	1984	2110	2311	2785	3039	3229	3070	2659	2543	2237	21
1965 10	1962	1910	2216	2437	2817	3123	3345	3112	2659	2469	2332	21
1966 05	1910	1941	2216	2342	2923	3229	3513	3355	2849	2680	2395	22
1967 47	1994	1952	2290	2395	2965	3239	3608	3524	3018	2648	2363	22
1968 84	1994	1941	2258	2332	3323	3608	3957	3672	3155	2933	2585	23
1969 61	2057	2100	2458	2638	3292	3724	4652	4379	4231	3756	3429	34
1970 88	3345	4220	4874	5064	5951	6774	7997	7523	7438	6879	6489	62
1971 54	5919	6183	6594	6489	8040	9715	9714	9756	8595	7861	7753	81
1972 29	7778	7402	8903	9742	11372	12741	13733	13691	12239	12502	11241	108
1973 53	11569	10397	12493	11962	13974	14945	16805	16587	14225	14157	13016	122
1974 16	11704	12275	13695	14082	16555	17339	17777	17592	16194	15336	14208	131
1975 12	12354	12682	14141	14989	16159	18276	19157	18737	17109	17094	15418	143
1976 98	13260	14990	15975	16770	19819	20983	22001	22337	20750	19969	17293	164
1977 88	15117	16058	18137	18471	21398	23854	26025	25479	22804	19619	19627	184
1978 95	17243	18284	20226	20903	23768	26323	28038	26776	22886	22813	22404	197
1979 30	18839	18892	20823	22212	25076	26884	30611	30228	26762	25885	23328	219
1980 48	21433	22369	24503	25905	30605	34984	37060	34502	31793	29275	28305	252
1981 45	27730	27424	32684	31366	37459	41060	43558	42398	33827	34962	33480	324
1982 29	30715	30400	31451	31306	40592	44133	47387	41310	37913	34355	34607	287
1983 34	26138	30745	35018	34549	40980	42869	45022	40387	38180	38608	35308	302
1984 14	28801	33034	35294	33181	40797	42355	46098	42430	41851	39331	37328	345
1985 42	32494	33308	36805	34221	41020	44350	46173	44435	40943	39269	35901	321
1986 79	31239	32261	34951	38109	43168	45547	49568	45387	41805	41281	36068	348
1987 58	32791	34206	39128	40249	43519	46137	56709	52306	49397	45500	39857	379
1988 95	35567	37696	42319	39137	47062	50610	54457	54435	48516	43225	42155	399

```

1989 37541 37277 41778 41666 49616 57793 61884 62400 50820 51116 45731 425
28
1990 40459 40295 44147 42697 52561 56572 56858 58363 45627 45622 41304 360
16
1991 35592 35677 39864 41761 50380 49129 55066 55671 49058 44503 42145 386
98
1992 38963 38690 39792 42545 50145 58164 59035 59408 55988 47321 42269 396
06
1993 37059 37963 31043 41712 50366 56977 56807 54634 51367 48073 46251 437
36
1994 39975 40478 46895 46147 55011 57799 62450 63896 57784 53231 50354 384
10
1995 41600 41471 46287 49013 56624 61739 66600 60054
> ### Checking the class of the imported ###
class(gasprod)
[1] "ts"
### Inspection of the time series data ###
summary(gasprod)
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
    1646    2675    16788    21415    38629    66600
anyNA(gasprod)
[1] FALSE
> findfrequency(gasprod)
[1] 12
> ts.plot(gasprod,gpars = list(xlab = "Year",ylab = "Gas Production",
+                             main = "Australian Gas Production (1956-199
+                             5)",
+                             col = c("Blue")),lwd = 2)
> stepplot = dygraph(gasprod,main = "Australian Gas Production (1956-1995)
+                    ,
+                    xlab = "Year",ylab = "Gas Production") %>%dyOptions(s
+                    tepPlot= TRUE,pointSize = 0,fillGraph = TRUE,
+                    fillAlpha = 0.2)
> stepplot
> ### Inspection of individual elements by decomposition of time series ##
##
> dc.gasprod = stl(gasprod,s.window = "periodic")
> plot(dc.gasprod)
> monthplot(gasprod,main = "Month Plot for Australian Gas Production")
> seasonplot(gasprod,year.labels = TRUE,
+            main = "Month Plot for Australian Gas Production")
> ### De-Seasonalizing the time series from the decomposed time series####
> ds.gasprod = (dc.gasprod$time.series[,2]+dc.gasprod$time.series[,3])
> ts.plot(ds.gasprod,gpars = list(xlab = "Year",ylab = "Gas Production",
+                                 main = "Australian Gas Production (Deseason
+                                 lized)",
+                                 col = c("Red")),lwd = 2)
> ts.plot(ds.gasprod,gasprod,gpars = list(xlab = "Year",ylab = "Gas Produc
+                                 tion",
+                                 main = "Australian Gas Production (Deseas
+                                 onlized
+                                 vs. Original)",
+                                 col = c("Red","Blue")),lwd = 2)
> ### Splitting the time series into training and testing samples (Origina
+ l) ####
> train.gasprod = window(gasprod,start=c(1970,1), end=c(1993,12), freq=12)
> ts.plot(train.gasprod)
> test.gasprod = window(gasprod,start=c(1994,1),end=c(1995,8), freq=12)
> ts.plot(test.gasprod)
> autoplot(train.gasprod, series="Train") + autolayer(test.gasprod, series
+ ="Test") +
+   ggtitle("Gas Produciton Traning and Test data") +
+   xlab("Year") + ylab("Production") +
+   guides(colour=guide_legend(title="Forecast"))
> ### Splitting the time series into training and testing samples(Deseason
+ alize) ####
> train.ds.gasprod = window(ds.gasprod,start=c(1970,1), end=c(1993,12), fr
+ eq=12)
> ts.plot(train.ds.gasprod)

```

```

> test.ds.gasprod = window(ds.gasprod,start=c(1994,1),end=c(1995,8), freq=
12)
> ts.plot(test.ds.gasprod)
> autoplot(train.ds.gasprod, series="Train") + autolayer(test.ds.gasprod,
series="Test") +
+ ggtitle("Gas Produciton Traning and Test data(Deseasonalized)") +
+ xlab("Year") + ylab("Production") +
+ guides(colour=guide_legend(title="Forecast"))
> ### Checking the periodicity of the Time Series ###
> periodicity(gasprod)
Monthly periodicity from Jan 1956 to Aug 1995
> findfrequency(gasprod)
[1] 12
> periodicity(train.gasprod)
Monthly periodicity from Jan 1970 to Dec 1993
> findfrequency(train.gasprod)
[1] 12
> periodicity(test.gasprod)
Monthly periodicity from Jan 1994 to Aug 1995
> findfrequency(test.gasprod)
[1] 12
> ##### Naive Method #####
> gasprod.rw = stl(train.gasprod,s.window = 'p')
> gasprod.rw = forecast(de.gasprod.rw,method = "rwdrift",h = 20)
> ts.plot(test.gasprod,gasprod.rw$mean,gpars = list(col = c("Red","Blue"),
+ main = "Random Walk with Drift(Original vs. Forecasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"
),lty = 1,
+ box.lwd = 0.1,cex = 0.75)
> vec.rw = cbind(test.gasprod,gasprod.rw$mean)
> MAPE.rw = mean(abs(vec.rw[,1]-vec.rw[,2])/vec.rw[,1])
> print(MAPE.rw)
[1] 0.0737419
> ##### Simple Exponential Smoothing (Original) #####
> gasprod.ses = ses(train.gasprod,start = c(1970,1),end = c(1993,12),frequ
ency = 12,h = 20)
> summary(gasprod.ses)

```

Forecast method: Simple exponential smoothing

Model Information:
Simple exponential smoothing

Call:
ses(y = train.gasprod, h = 20, start = c(1970, 1), end = c(1993,

call:
12), frequency = 12)

Smoothing parameters:
alpha = 0.9999

Initial states:
l = 6002.9039

sigma: 3344.025

	AIC	AICc	BIC
	6309.126	6309.210	6320.115

Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE
ACF1						
Training set	131.0317	3332.394	2416.021	0.1509165	8.113316	0.9164853
0855						

Forecasts:

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Jan 1994	43736.25	39450.71	48021.79	37182.08	50290.42


```
> print(gasprod.ses$mean)
```

	Nov	Dec
1994	43736.25	43736.25
1995		

Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE
ACF1						
Training set	139.963	2311.279	1591.754	0.2300846	6.077039	0.6038107
553574						-0.005

```

Forecasts:
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+       box.lwd = 0.1,cex = 0.75)
> vec.ses = cbind(test.gasprod,gasprod.ses$mean)
> MAPE.ses = mape(test.gasprod,gasprod.ses$mean)
> print(MAPE.ses)
[1] 0.1726068
> #### Simple Exponential Smoothing (Deseasonalized) ####
> gasprod.dses = ses(train.ds.gasprod,start = c(1970,1),end = c(1993,12),f
frequency = 12,h = 20)
> summary(gasprod.dses)

```

Forecast method: Simple exponential smoothing

Model Information:
Simple exponential smoothing

Call:
ses(y = train.ds.gasprod, h = 20, start = c(1970, 1), end = c(1993,

Call:
12), frequency = 12)

Smoothing parameters:
alpha = 0.9833

Initial states:
l = 7432.9827

sigma: 2319.347

	AIC	AICc	BIC
6098.373	6098.458	6109.362	

Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE
ACF1						
Training set	139.963	2311.279	1591.754	0.2300846	6.077039	0.6038107
553574						-0.005

```

Forecasts:
      Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
Jan 1994      47067.77 44095.41 50040.13 42521.93 51613.60
Feb 1994      47067.77 42899.24 51236.30 40692.55 53442.99
Mar 1994      47067.77 41976.76 52158.78 39281.74 54853.80
Apr 1994      47067.77 41197.50 52938.04 38089.96 56045.58
May 1994      47067.77 40510.19 53625.35 37038.82 57096.72
Jun 1994      47067.77 39888.39 54247.15 36087.85 58047.69
Jul 1994      47067.77 39316.30 54819.24 35212.92 58922.62
Aug 1994      47067.77 38783.63 55351.91 34398.27 59737.27
Sep 1994      47067.77 38283.20 55852.34 33632.93 60502.61
Oct 1994      47067.77 37809.78 56325.76 32908.89 61226.65
Nov 1994      47067.77 37359.42 56776.12 32220.12 61915.41
Dec 1994      47067.77 36929.04 57206.50 31561.92 62573.62
Jan 1995      47067.77 36516.20 57619.34 30930.54 63205.00
Feb 1995      47067.77 36118.92 58016.62 30322.95 63812.59
Mar 1995      47067.77 35735.56 58399.98 29736.65 64398.89
Apr 1995      47067.77 35364.75 58770.79 29169.54 64966.00
May 1995      47067.77 35005.33 59130.21 28619.86 65515.68
Jun 1995      47067.77 34656.32 59479.22 28086.09 66049.45
Jul 1995      47067.77 34316.85 59818.69 27566.92 66568.61
Aug 1995      47067.77 33986.19 60149.34 27061.23 67074.31
> print(gasprod.dses$mean)

```

	Jan	Feb	Mar	Apr	May	Jun	Jul	A
1994	47067.77	47067.77	47067.77	47067.77	47067.77	47067.77	47067.77	47067.77
1995	47067.77	47067.77	47067.77	47067.77	47067.77	47067.77	47067.77	47067.77

	Nov	Dec
1994	47067.77	47067.77
1995		

```

> ts.plot(test.ds.gasprod,gasprod.dses$mean,gpars = list(col = c("Red","Blue"),
+ main = "Simple Exponential Smoothing(Original vs. Forecasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+ box.lwd = 0.1,cex = 0.75)
> MAPE.dses = mape(test.ds.gasprod,gasprod.dses$mean)
> print(MAPE.dses)
[1] 0.1110184
> ##### Double Exponential Method (Holt Model) (Original) #####
> gasprod.holt = holt(train.gasprod ,start=c(1970,1),end=c(1993,12), freq=12,h=20)
> summary(gasprod.holt)

```

Forecast method: Holt's method

Model Information:
Holt's method

Call:
holt(y = train.gasprod, h = 20, start = c(1970, 1), end = c(1993,

call:
12), freq = 12)

Smoothing parameters:
alpha = 0.9436
beta = 0.5881

Initial states:
l = 3775.0087
b = 467.2663

sigma: 3542.179

	AIC	AICC	BIC
	6344.263	6344.476	6362.578

Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE
ACF1						
Training set	-16.95919	3517.494	2436.963	0.2550371	7.909274	0.9244294
1120745						-0.0

Forecasts:

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Jan 1994	41337.3687	36797.884	45876.85	34394.826	48279.91
Feb 1994	38932.0524	30628.005	47236.10	26232.108	51632.00
Mar 1994	36526.7362	23815.946	49237.53	17087.260	55966.21
Apr 1994	34121.4199	16438.554	51804.29	7077.810	61165.03
May 1994	31716.1037	8552.086	54880.12	-3710.204	67142.41
Jun 1994	29310.7874	199.210	58422.36	-15211.528	73833.10
Jul 1994	26905.4711	-8586.447	62397.39	-27374.733	81185.68
Aug 1994	24500.1549	-17777.552	66777.86	-40158.018	89158.33
Sep 1994	22094.8386	-27351.339	71541.02	-53526.565	97716.24
Oct 1994	19689.5224	-37288.475	76667.52	-67450.805	106829.85
Nov 1994	17284.2061	-47572.276	82140.69	-81905.224	116473.64
Dec 1994	14878.8899	-58188.160	87945.94	-96867.519	126625.30
Jan 1995	12473.5736	-69123.236	94070.38	-112317.977	137265.12

```

Feb 1995      10068.2573  -80366.005  100502.52  -128239.012  148375.53
Mar 1995      7662.9411  -91906.127  107232.01  -144614.807  159940.69
Apr 1995      5257.6248  -103734.234  114249.48  -161431.039  171946.29
May 1995      2852.3086  -115841.794  121546.41  -178674.656  184379.27
Jun 1995      446.9923   -128220.987  129114.97  -196333.701  197227.69
Jul 1995     -1958.3239  -140864.616  136947.97  -214397.166  210480.52
Aug 1995     -4363.6402  -153766.027  145038.75  -232854.874  224127.59
> gasprod.holt$mean
      Jan      Feb      Mar      Apr      May      Jun
Jul 1994 41337.3687 38932.0524 36526.7362 34121.4199 31716.1037 29310.7874 269
05.4711 24500.1549
1995 12473.5736 10068.2573 7662.9411 5257.6248 2852.3086 446.9923 -19
58.3239 -4363.6402
      Sep      Oct      Nov      Dec
1994 22094.8386 19689.5224 17284.2061 14878.8899
1995
> ts.plot(test.gasprod,gasprod.holt$mean,gpars = list(col = c("Red","Blue"
),
+
+ main = "Holt's Method(
Original vs. Forecasted)")
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"
),lty = 1,
+ box.lwd = 0.1,cex = 0.75)
> MAPE.holt = mape(test.gasprod,gasprod.holt$mean)
> print(MAPE.holt)
[1] 0.6201079
> #### Double Exponential Method (Holt Model) (Deseasonlaize) ####
> gasprod.dholt = holt(train.ds.gasprod,start = c(1970,1),end = c(1993,12)
,freq = 12,h = 20)
> summary(gasprod.dholt)

```

Forecast method: Holt's method

Model Information:
Holt's method

```

Call:
holt(y = train.ds.gasprod, h = 20, start = c(1970, 1), end = c(1993,
12), freq = 12)

Smoothing parameters:
alpha = 0.9809
beta = 1e-04

Initial states:
l = 7297.2832
b = 70.1218

sigma: 2324.336

      AIC      AICC      BIC
6101.590 6101.803 6119.905

```

```

Error measures:
      ME      RMSE      MAE      MPE      MAPE      MASE
ACF1
Training set 67.98246 2308.139 1589.453 -0.1147602 6.078146 0.6029375 -0.0
02865917

```

```

Forecasts:
      Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
Jan 1994      47142.87 44164.12 50121.63 42587.26 51698.49
Feb 1994      47214.95 43042.25 51387.66 40833.36 53596.55
Mar 1994      47287.03 42192.76 52381.31 39496.01 55078.06
Apr 1994      47359.11 41485.99 53232.24 38376.95 56341.28
May 1994      47431.19 40870.91 53991.47 37398.11 57464.27

```

Jun 1994	47503.27	40321.17	54685.38	36519.19	58487.36
Jul 1994	47575.35	39821.01	55329.70	35716.11	59434.60
Aug 1994	47647.43	39360.17	55934.70	34973.16	60321.71
Sep 1994	47719.51	38931.49	56507.54	34279.39	61159.64
Oct 1994	47791.59	38529.75	57053.44	33626.82	61956.36
Nov 1994	47863.67	38151.00	57576.34	33009.43	62717.92
Dec 1994	47935.75	37792.19	58079.31	32422.52	63448.99
Jan 1995	48007.83	37450.87	58564.79	31862.36	64153.31
Feb 1995	48079.91	37125.06	59034.76	31325.91	64833.91
Mar 1995	48151.99	36813.13	59490.86	30810.69	65493.29
Apr 1995	48224.07	36513.70	59934.44	30314.61	66133.53
May 1995	48296.15	36225.64	60366.67	29835.89	66756.41
Jun 1995	48368.23	35947.94	60788.52	29373.04	67363.43
Jul 1995	48440.31	35679.76	61200.86	28924.73	67955.89
Aug 1995	48512.39	35420.35	61604.43	28489.84	68534.94

```
> gasprod.dholt$mean
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug
1994	47142.87	47214.95	47287.03	47359.11	47431.19	47503.27	47575.35	47647.43
1995	48007.83	48079.91	48151.99	48224.07	48296.15	48368.23	48440.31	48512.39

	Nov	Dec
1994	47863.67	47935.75
1995		

```
> ts.plot(test.ds.gasprod,gasprod.dholt$mean,gpars = list(col = c("Red","Blue"),
+
+ main = "Holt's Method(Original vs. Forecasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+ box.lwd = 0.1,cex = 0.75)
> MAPE.dholt = mape(test.ds.gasprod,gasprod.dholt$mean)
> print(MAPE.dholt)
[1] 0.1033223
> ##### Holt winter's method (Original) #####
> gasprod.hw = hw(train.gasprod,start = c(1970,1),end = c(1993,12),freq = 12,h = 20)
> summary(gasprod.hw)
```

Forecast method: Holt-Winters' additive method

Model Information:
Holt-Winters' additive method

Call:
hw(y = train.gasprod, h = 20, start = c(1970, 1), end = c(1993,

call:
12), freq = 12)

Smoothing parameters:
alpha = 0.3408
beta = 1e-04
gamma = 0.5936

Initial states:
l = 6253.203
b = 119.5506
s = -4511.742 -2141.073 234.0438 2010.202 5919.329 7284.465
5272.426 2485.985 -2602.642 -3068.389 -5131.983 -5750.62

sigma: 2109.356

AIC	AICc	BIC
6057.255	6059.521	6119.525

Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE
ACF1						
Training set	78.69136	2049.926	1551.842	0.4623734	7.505829	0.5886704
8151						0.273

Forecasts:

	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Jan 1994		40279.65	37576.40	42982.90	36145.38	44413.91
Feb 1994		41289.25	38433.25	44145.25	36921.37	45657.13
Mar 1994		38304.76	35303.68	41305.84	33715.01	42894.51
Apr 1994		47712.40	44572.87	50851.93	42910.90	52513.90
May 1994		55810.29	52538.08	59082.50	50805.88	60814.70
Jun 1994		61602.40	58202.61	65002.18	56402.87	66801.92
Jul 1994		61619.91	58097.10	65142.73	56232.23	67007.60
Aug 1994		60595.61	56953.85	64237.37	55026.02	66165.20
Sep 1994		57042.03	53285.03	60799.04	51296.18	62787.88
Oct 1994		51962.18	48093.29	55831.07	46045.23	57879.14
Nov 1994		48460.99	44483.30	52438.68	42377.64	54544.34
Dec 1994		44992.38	40908.73	49076.03	38746.98	51237.78
Jan 1995		41741.45	36938.07	46544.84	34395.31	49087.59
Feb 1995		42751.05	37859.46	47642.65	35270.01	50232.10
Mar 1995		39766.56	34788.27	44744.85	32152.93	47380.20
Apr 1995		49174.20	44110.65	54237.76	41430.17	56918.24
May 1995		57272.09	52124.64	62419.54	49399.75	65144.44
Jun 1995		63064.20	57834.15	68294.25	55065.53	71062.88
Jul 1995		63081.72	57770.30	68393.13	54958.61	71204.83
Aug 1995		62057.41	56665.82	67449.01	53811.68	70303.15

> gasprod.hw\$mean

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug
1994	40279.65	41289.25	38304.76	47712.40	55810.29	61602.40	61619.91	60595.61
1995	41741.45	42751.05	39766.56	49174.20	57272.09	63064.20	63081.72	62057.41

```

> ts.plot(test.gasprod,gasprod.hw$mean,gpars = list(col = c("Red","Blue"),
+                                                    main = "Holt Winter's(
Original vs. Forecasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"
),lty = 1,
+       box.lwd = 0.1,cex = 0.75)
> MAPE.hw = mape(test.gasprod,gasprod.hw$mean)
> print(MAPE.hw)
[1] 0.04666037
> ##### Holt winter's method (Deseasonalize) #####
> gasprod.dhw = hw(train.ds.gasprod,start = c(1970,1),end = c(1993,12),fre
q = 12,h = 20)
> summary(gasprod.dhw)

```

Forecast method: Holt-winters' additive method

Model Information:
Holt-winters' additive method

Call:
hw(y = train.ds.gasprod, h = 20, start = c(1970, 1), end = c(1993,

12), freq = 12)

Smoothing parameters:
alpha = 0.3408
beta = 1e-04
gamma = 0.5934

Initial states:
l = 6263.4243

```

b = 119.8431
s = -1190.992 -669.5463 69.8913 481.8762 1663.536 1972.405
    1680.539 771.3246 -793.5156 -903.362 -1416.127 -1666.028

```

```
sigma: 2109.394
```

```

      AIC      AICC      BIC
6057.265 6059.531 6119.535

```

```
Error measures:
```

```

      ME      RMSE      MAE      MPE      MAPE      MASE
ACF1
Training set 77.89129 2049.962 1551.948 -1.123422 8.141496 0.5887106 0.273
9442

```

```
Forecasts:
```

```

      Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
Jan 1994      44364.14 41660.84 47067.44 40229.80 48498.47
Feb 1994      45005.85 42149.82 47861.88 40637.93 49373.77
Mar 1994      40472.48 37471.40 43473.56 35882.73 45062.24
Apr 1994      49522.52 46383.00 52662.03 44721.04 54323.99
May 1994      54097.59 50825.41 57369.76 49093.23 59101.95
Jun 1994      58011.54 54611.81 61411.28 52812.09 63210.99
Jul 1994      56309.27 52786.52 59832.02 50921.68 61696.85
Aug 1994      56342.08 52700.40 59983.76 50772.61 61911.55
Sep 1994      55515.56 51758.64 59272.48 49769.85 61261.27
Oct 1994      51799.98 47931.20 55668.77 45883.19 57716.78
Nov 1994      49934.90 45957.33 53912.47 43851.73 56018.07
Dec 1994      48315.63 44232.11 52399.15 42070.43 54560.84
Jan 1995      45829.18 41026.13 50632.22 38483.56 53174.79
Feb 1995      46470.89 41579.64 51362.13 38990.37 53951.40
Mar 1995      41937.52 36959.58 46915.45 34324.42 49550.61
Apr 1995      50987.55 45924.36 56050.74 43244.07 58731.04
May 1995      55562.62 50415.54 60709.71 47690.83 63434.41
Jun 1995      59476.58 54246.90 64706.26 51478.47 67474.69
Jul 1995      57774.30 52463.27 63085.34 49651.77 65896.84
Aug 1995      57807.12 52415.90 63198.33 49561.97 66052.27

```

```
> gasprod.dhw$mean
```

```

      Jan      Feb      Mar      Apr      May      Jun      Jul      A
ug      Sep      Oct
1994 44364.14 45005.85 40472.48 49522.52 54097.59 58011.54 56309.27 56342.
08 55515.56 51799.98
1995 45829.18 46470.89 41937.52 50987.55 55562.62 59476.58 57774.30 57807.
12
      Nov      Dec
1994 49934.90 48315.63
1995

```

```

> ts.plot(test.ds.gasprod,gasprod.dhw$mean,gpars = list(col = c("Red","Blue"),
+
+                                     main = "Holt winter's(
Original vs. Forecasted)"))
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+       box.lwd = 0.1,cex = 0.75)
> MAPE.dhw = mape(test.ds.gasprod,gasprod.dhw$mean)
> print(MAPE.dhw)
[1] 0.0458347
> ##### Checking for stationarity using visualization #####
> ts.plot(gasprod,gpars = list(xlab = "Year",ylab = "Gas Production",
+
+                               main = "Australian Gas Production (1960-1995)",
+
+                               col = c("Blue"))),lwd = 2)
> ##### Plotting ACF and PACF plots #####
> acf(gasprod, lag.max = 50)
> pacf(gasprod, lag.max = 50)
> ##### Checking for stationarity using Augmented Dicky- Fuller Test #####
> adf = adf.test(gasprod,alternative = "stationary")
> adf

```

Augmented Dickey-Fuller Test

data: gasprod
Dickey-Fuller = -2.7131, Lag order = 7, p-value = 0.2764
alternative hypothesis: stationary

```
> print(adf$alternative)
[1] "stationary"
> if(adf$p.value < 0.05){
+   print("The series is Stationary & Null Hypothesis is rejected")
+ } else{
+   print("The Series is not Stationary & Null Hypothesis is accepted")
+ }
[1] "The Series is not Stationary & Null Hypothesis is accepted"
> #### Stationarizing the series for performing Manual Arima ####
> diff.gasprod = diff(gasprod, differences = 1)
> ts.plot(diff.gasprod,col = c("Blue"),lwd = 1,main = "Differenced Series"
)
> abline(a=1,b=0,col = c("Red"),lwd =4)
> acf(diff.gasprod, lag.max = 50)
> pacf(diff.gasprod, lag.max = 50)
> adf.d = adf.test(diff.gasprod,alternative = "stationary")
Warning message:
In adf.test(diff.gasprod, alternative = "stationary") :
  p-value smaller than printed p-value
> adf.d
```

Augmented Dickey-Fuller Test

data: diff.gasprod
Dickey-Fuller = -19.321, Lag order = 7, p-value = 0.01
alternative hypothesis: stationary

```
> print(adf.d$alternative)
[1] "stationary"
> if(adf.d$p.value < 0.05){
+   print("The series is Stationary & Null Hypothesis is rejected")
+ } else{
+   print("The Series is not Stationary & Null Hypothesis is accepted")
+ }
[1] "The series is Stationary & Null Hypothesis is rejected"
> ### Splitting of time series into Training and Testing Data ###
> diff.train.gasprod = window(diff.gasprod,start = c(1970,1),end = c(1993,
12),frequency = 12)
> diff.test.gasprod = window(diff.gasprod,start = c(1994,1),end = c(1995,8
),frequency = 12)
> plot.ts(diff.train.gasprod,main = "Differenced Train")
> plot.ts(diff.test.gasprod,main = "Differenced Test")
> ### Finding the P and Q values ###
> acf(diff.train.gasprod)
> pacf(diff.train.gasprod)
> ### Building the Manual ARIMA model ###
> marima.gasprod = Arima(diff.train.gasprod,order = c(2,1,2))
> marima.gasprod
Series: diff.train.gasprod
ARIMA(2,1,2)

Coefficients:
      ar1      ar2      ma1      ma2
    -0.0131  0.2545  -0.7212  -0.2788
s.e.    0.1623  0.0695   0.1595   0.1593

sigma^2 estimated as 9895599:  log likelihood=-2719.06
AIC=5448.12   AICc=5448.33   BIC=5466.41
> forc.marima.gasprod = forecast(marima.gasprod,h = 20)
> ts.plot(diff.test.gasprod,forc.marima.gasprod$mean,gpars = list(col = c(
"Red","Blue"),
+                                                                    main = "Manual Arima(O
riginal vs. Forecasted)"))
```



```

> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+       box.lwd = 0.1,cex = 0.50)
> vec.marima = cbind(diff.test.gasprod,forc.marima.gasprod$mean)
> MAPE.marima = mean(abs(vec.marima[,1]-vec.marima[,2])/vec.marima[,1])
> print(MAPE.marima)
[1] 0.1915037
> ###USING AUTO ARIMA (Original) ###
> aarima.gasprod = auto.arima(train.gasprod,seasonal = TRUE)
> aarima.gasprod
Series: train.gasprod
ARIMA(2,1,1)(0,1,2)[12]

Coefficients:
          ar1      ar2      ma1      sma1      sma2
      0.5017  0.2057 -0.9583 -0.4404 -0.1236
s.e.  0.0738  0.0722  0.0426  0.0676  0.0639

sigma^2 estimated as 3535010:  log likelihood=-2463.67
AIC=4939.33  AICc=4939.64  BIC=4961.03
> forc.aarima = forecast(aarima.gasprod,h = 20)
> ts.plot(test.gasprod,forc.aarima$mean,gpars = list(col = c("Red","Blue"),
+       box.lwd = 0.1,cex = 0.75))
main = "Auto Arima(Original vs. Forecasted)")
> legend("topleft", legend = c("Actual","Forecasted"),col = c("Red","Blue"),lty = 1,
+       box.lwd = 0.1,cex = 0.75)
> vec.aarima = cbind(test.gasprod,forc.aarima$mean)
> MAPE.aarima = mean(abs(vec.aarima[,1]-vec.aarima[,2])/vec.aarima[,1])
> print(MAPE.aarima)
[1] 0.06370687
> ### Ljung box test ###
> Box.test(aarima.gasprod$residuals,type = "Ljung-Box")

```

Box-Ljung test

```

data: aarima.gasprod$residuals
X-squared = 0.009919, df = 1, p-value = 0.9207

> hist(aarima.gasprod$residuals,main = "Histogram of Auto ARIMA residuals",
+       xlab = "Residuals")
> ### Making the future forecast using the best model ###
> future = forecast(aarima.gasprod,h = 32)
> plot(future)
> future$mean
      Jan      Feb      Mar      Apr      May      Jun      Jul      A
1994 40911.17 41136.45 38318.03 44418.69 52575.33 57827.42 59054.39 57800.
39 52923.33 49148.06
1995 41314.97 41577.63 40068.39 45141.64 53284.34 58496.89 60063.56 59142.
29 54005.17 49942.73
1996 41882.79 42199.58 40744.41 45855.92 54028.93 59264.56 60849.03 59941.
44
      Nov      Dec
1994 46382.99 43491.75
1995 46875.98 43928.71
1996

> ### Making the future forecast using the model created from gasprod ###
> aarima.fgasprod = auto.arima(gasprod,seasonal = TRUE)
> fgasprod.forc = forecast(aarima.fgasprod,h = 12)
> plot(fgasprod.forc)
> fgasprod.forc$mean
      Jan      Feb      Mar      Apr      May      Jun      Jul      A
1995 56907.83 52476.28

```

1996	43318.41	43601.71	46668.11	49376.36	57536.25	62184.69	65795.74	63391.54
		Nov						
1995	49719.79		43473.70					
1996								